

RAHUL SHETTY

United States |  Email |  Github |  Portfolio | +1 (720) 803-9791

EDUCATION

- University of Colorado, Boulder** - Master of Science, Computer Science | GPA: 3.84/4.0 Aug 2024 – May 2026
- **Coursework:** Machine Learning, Natural Language Processing, Datacenter Scale Computing, Distributed Systems
- Sardar Patel Institute of Technology** - Bachelor of Technology, ECE | GPA: 9.1/10 Aug 2018 – May 2022
- **Coursework:** Data Structures and Algorithms, Programming (Python, Java), Machine Learning

SKILLS

Languages: Python, Java, SQL, MATLAB, BASH scripting, JavaScript, HTML5, CSS3

ML & AI: TensorFlow, PyTorch, Pandas, Scikit-learn, Matplotlib, GenAI, CNN, Transformers, RAG, Knowledge Graphs, VLA

Frameworks & Libraries: React JS, ROS2, NumPy, LangChain, Hadoop, Spark, Streamlit

Technologies: Ansible, Docker, GitHub, Google Cloud Platform, Linux, AWS, ChromaDB, Pinecone, Kubernetes, Firebase

EXPERIENCE

- Graduate Researcher, Human Interaction and Robotics Group – CU Boulder** Jan 2025 – May 2026
- Developed **DEFT, a diffusion-based AI framework** for robot resilience that improved task success from **42% → 99%** and constraint satisfaction from **37% → 75%**, evaluated across **4.7M+ trajectories** and **2,400+ failure scenarios**.
 - Engineered **task-conditioned diffusion models (using FiLM)** to enable zero-shot adaptation to **unseen failures** without retraining, achieving **74% success on out-of-distribution errors** and demonstrating robust autonomy in **real-world hardware tests**.
 - Manuscript accepted at **IEEE International Conference on Robotics and Automation (ICRA)** (arXiv:2602.02895).

- Software Development Engineer, Nomura** July 2022 – July 2024
- Enhanced the Front Office Supervision (FOS) portal by creating an **organizational hierarchy tree** from data fetched from the **Oracle** database and refactored parts of the **SpringBoot** code.
 - Collaborated with a team to migrate **700+ data marts (13TB)** from **Sybase** to **Snowflake** on **AWS**, saving over **\$800K** in **infrastructure costs** while modernizing **three large-scale applications**.
 - Contributed to the development of a **Snowpark** framework using **Python**, **Shell scripting** and **SnowSQL** to streamline **Informatica** workflows and cut data processing time by **75%**.

- Software Engineering Intern, Nomura** Jan 2022 – June 2022
- **Automated** retrieval of **Git** details for production releases through **JIRA API integration**, streamlining release preparation and **eliminating human errors**. Developed a scalable **Python** backend with a **Django** web application frontend, supporting **10+ teams**.
 - **Led the migration** of an existing JSP-based monitoring portal to a **React.js** frontend and a refactored **SpringBoot** backend, resulting in a **100%** code overhaul.

PROJECTS

- Agentic Hybrid RAG Engine |  | Python, Neo4j, Qdrant, LangChain, Docker, Streamlit** Dec 2025 - Jan 2026
- Built a **production-grade Hybrid Retrieval Engine** that routes queries between **Structured (Graph)** and **Unstructured (Vector)** stores, **optimizing latency and relevance** for enterprise use cases.
 - Implemented **Agentic Memory** to personalize responses based on **user roles (CTO vs. CEO)**, adapting technical depth and **boosting relevance scores** by **25%** in **A/B testing** scenarios.
 - Benchmarked system reliability against **4 industry baselines (including BM25 & HyDE)**, demonstrating **3x higher recall** on exhaustive list generation and **100% precision** on fact-verification tasks.

- Photo Memory Finder |  | Python, GCP (Vertex AI, Cloud Run), React, Pinecone** Nov 2025 – Dec 2025
- Architected a **scalable distributed** event-driven pipeline on **Google Cloud Platform** capable of processing concurrent photo uploads with **<2s latency**, achieving semantic understanding beyond simple metadata.
 - Built a **multimodal** ingestion pipeline using **Cloud Pub/Sub** and **Cloud Run** workers to asynchronously process uploads, reducing metadata extraction time by **40%** via parallelized **hybrid embedding generation (image + text)** with **Vertex AI**.
 - Implemented an intelligent LLM Reranking system using **Gemini 2.5 Flash**, filtering and re-ordering vector search results based on query context, improving **retrieval precision by 35%** for complex queries.

- Ansh.ai – Real-time Conversational AI Journal |  | Python, Flask, React, LangChain, LLM** Feb 2025 - Mar 2025
- Built an AI-powered journaling companion using **Python (Flask)** and **React**, integrating **LangChain** and **ChromaDB**, and leveraging **Retrieval-Augmented Generation (RAG)** to enhance conversations with relevant context from past user interactions.
 - Deployed on **AMD Minisforum AI 370**, leveraging its Hybrid NPU + GPU for efficient on-device LLM inference and **multimodal** interaction through a talking AI avatar.
 - Achieved **2nd place** in the **AMD AI track at HackCU 11**, recognized for innovation in real-time, resource-efficient AI applications.